

Frozen element embeddings versus categorical species in linear ACE: distillation benchmarks on CrMnFeCoNi

ACEpotentials.jl jax-eval branch — 14 September 2026

Summary

A linear ACE model whose species dependence enters through a **frozen element embedding** (Darby-style tensor reduction, channel folded into the radial index) was benchmarked against the standard **categorical** model on a five-component alloy, with labels distilled from MACE-MH-1. At matched polynomial degree the embedded model costs 7–17% on force RMSE for 3–6× fewer parameters; given one more degree — which its smaller design matrix makes affordable — it **matches** the categorical model at 2.4× fewer parameters (0.0767 vs 0.0763 eV/Å). The best absolute result is categorical degree 10 on 4 000 structures, 0.0641 eV/Å and 2.0 meV/atom, fitted from a 151 GB design matrix. Learning curves show the categorical model is data-limited and the embedded model basis-limited, which is why the trade runs the way it does. Two surprises: the embedding’s **values** do not matter at $S = 5$, and the Julia evaluator ran the embedded model **slower** than the categorical one — now profiled and fixed. Recommendations, including code changes in flight, close the report.

Set-up

System and labels. Rattled and strained fcc CrMnFeCoNi cells (32–48 atoms, random equiatomic occupancy, no short-range order), generated with bounded uniform perturbations (lattice $\pm 3\%$, shear 0.02, rattle 0.02–0.10 Å, minimum separation 2.0 Å). Energies, forces and virials from MACE-MH-1 on an RTX A4500 (4 000 structures in 4.6 min). An earlier 6%-Gaussian generator produced crushed cells (36% compression, forces above 20 eV/Å) carrying 77% of the squared test error; bounding it halved the error at fixed basis and budget. MACE-MP-0-small was rejected as teacher: a difference-fit showed its disagreement with MH-1 to be rough, not a smooth offset.

Models. Order 3, ACE1 heuristics, uniform cutoff 6.25 Å. Categorical `ace1_model` with $S = 5$ elements; embedded `ace_embedding_model` with per-order widths $d_\nu = \min(d_{\max}, \binom{S+\nu-1}{\nu})$: **lossless** [5, 15, 35] (spans the full species tensor at every order — a reparameterisation), $d_{\max} = 16$ ([5, 15, 16], order 3 cut to 16 of 35 directions) and $d_{\max} = 8$. Embedding table: MACE-MH-1’s 512-channel `node_embedding` reduced by PCA of the five rows (Gram preserved exactly; generic orthonormal mixing for $d > S$). Parameters are $n_B \times S$ (a readout per central species).

Fitting. Held-out 80/20 split; algebraic smoothness prior $p = 4$; weights from `acefit!` defaults; factor-once Tikhonov (QR of A , SVD of R) with λ swept per model and the best held-out force RMSE reported. Assembly on 12 workers; the degree-10 categorical matrix (404 144 × 46 885, 151 GB) was factorised in place on a 376 GB node.

Results

model	n_B	params	1k test F	4k test F	4k test E	rel. to cat.
deg 6 categorical	1 348	6 740	0.0969	—	—	—
deg 6 embedded $d \leq 16$	323	1 615	0.1101	—	—	+14%, 4.2× fewer (1k)
deg 8 categorical	3 824	19 120	0.0855	0.0763	0.0024	—
deg 8 embedded lossless	1 190	5 950	0.0906	0.0885	0.0025	+16%, 3.2× fewer
deg 8 embedded $d \leq 16$	753	3 765	0.0911	0.0894	0.0025	+17%, 5.1× fewer
deg 8 embedded $d \leq 8$	408	2 040	0.0986	0.1000	0.0027	+31%, 9.4× fewer
deg 10 categorical	9 327	46 635	—	0.0641	0.0020	—
deg 10 embedded lossless	2 730	13 650	—	0.0753	0.0022	+17%, 3.4× fewer
deg 10 embedded $d \leq 16$	1 609	8 045	—	0.0767	0.0022	+20%, 5.8× fewer
deg 10 embedded $d \leq 8$	850	4 250	—	0.0895	0.0025	+40%, 11× fewer

Table 1: Held-out force RMSE (eV/Å) and energy RMSE (eV/atom) against MH-1 labels; 1k = 800/200 split, 4k = 3 200/800. “rel. to cat.” compares with the categorical model at the same degree on 4k.

Three things follow from Table 1 and Figure 1. First, at matched degree the embedding costs 14–20% on forces (energies are unaffected: 2.2–2.7 meV/atom throughout), for 3–6× fewer parameters. Second, **lossless and $d \leq 16$ are indistinguishable at every degree** — truncating order 3 from 35 to 16 species directions costs nothing measurable at $S = 5$ — whereas $d \leq 8$, which also truncates order 2, costs about one degree. Third, per basis function the embedded model wins at every point: degree-10 $d \leq 16$ matches degree-8 categorical at 2.4× fewer parameters, and it was affordable (26 GB) where the categorical alternative needed a 376 GB node.




Figure 1: Test force RMSE against parameter count. Point labels are the polynomial degree. On the 4k set the degree-10 embedded $d \leq 16$ model (8 045 parameters) sits at the accuracy of degree-8 categorical (19 120).




Figure 2: Learning curves on the 1k set (test solid, train dashed). Degree 6 is basis-limited for both models; at degree 8 the categorical model is still data-limited (gap 0.022 at $N = 800$) while the embedded one has converged.

The learning curves (Figure 2) explain the table. Both degree-6 models saturate — train and test meet — by 400 structures; the embedded model gets there by 200 and is the better model below that. At degree 8 the categorical model still has slope at 800 and duly gains 11% from 4× the data, while the embedded models gain 2%. **The embedded basis wants degree, not data**; the categorical basis wants data. This is the mechanism behind “accurate per structure vs accurate per parameter”.

The embedding’s values do not matter at $S = 5$

A 2×2 of embedding source (MP-0-small, MH-1) × reduction (first- d column truncation, PCA) at degree 6 gave test F within 0.001 eV/Å at $d \leq 16$ and 0.005 at $d \leq 8$, although the truncated tables’ element-similarity structure correlates only 0.5–0.7 with the full tables’. At $S = 5$ the frozen embedding acts as a generic tensor reduction, consistent with Darby et al.’s random-projection result; transferred chemistry can only show once $d_{\max} < S$ truncates order 1 itself. PCA remains the right construction (it was also what exposed that a DCT mixing frame spans only 9/15 and 13/35 of the species tensors — a genericity check now in the tests).

Evaluation cost

The JAX evaluator’s GPU kernel (fp32) runs 32–38× MACE-MH-1’s at 40–1 300 atoms, 6–11× end-to-end with the host neighbour list; MH-1 does not fit a 3 000-atom cell on 20 GB. The Julia evaluator, however, ran the **embedded** model 1.8× slower than the categorical one despite 4× fewer basis functions, and both far below a well-written CPU kernel. Profiling found the cause: the splined radial tables are LEN wide (224 columns per edge for the embedded model), yet because the embedding weights make every R_{nl} a scalar multiple of $P_{n'}$, only 6–8 columns are distinct — the evaluator computes them all, as $\text{SVector}\{\text{LEN}, \text{Dual}\}$. A prototype exploiting the factorisation ($R_{nl} = c_{n,\text{pair}} P_{n'}$,




Figure 3: Throughput of the degree-6 $d \leq 16$ student on one host (lestrade: 32 cores, RTX 4000 Ada), 256–384-atom cells. The JAX kernel runs 32–38× the teacher; a single Julia core matches the teacher on a GPU.

the structure acejax already uses) with an allocation-free site loop is exact to 10^{-12} and 12–18× faster on four cores, above the JAX fp64 GPU kernel.

Recommendations

1. **Use the embedded model where parameters, memory or many elements are the constraint, and give it one more degree than the categorical model would get.** Degree-10 $d \leq 16$ is the current best trade (0.0767 eV/Å, 8 045 parameters, 26 GB matrix on 4k). Where accuracy per structure is all that matters and a large node is available, categorical degree 10 (0.0641) still wins; categorical degree 12 (330 GB) needs the distributed assembly of Stage 2 phase 18, embedded degree 12 (≈ 50 GB) does not.
2. **Keep $d_{\max} = 16$; do not use $d_{\max} = 8$** (it truncates order 2 and costs a full degree).
3. **Treat the embedding source as immaterial at $S \leq d_{\max}$** , and test the transferred-chemistry claim only at $S = 10$ –25, where the GRACE-FS distillation on HEA25 is the comparison to beat.
4. **Code changes in flight** (branch fix/basis-ed-performance off main, one PR). (a) Design-matrix path: evaluate_basis_ed replaced by a type-stable pushforward Jacobian and the boxed accumulation in energy_forces_virial_basis removed – 174–218× per structure, exact to 10^{-12} , 514 existing + 158 new tests passing; the 151 GB degree-10 assembly took 25 min instead of a projected 14 h. (b) Forward path: factorised radial splines with a plain-Float64 kernel, allocation-free site loop with $w_{AA} = A2B^T W_B$ folding, neighbour-list reuse – in progress on the same branch. Follow-ups elsewhere: ACEfit’s per-task model serialisation, GC.gc() per task and Array(A) copy; EquivariantTensors’ _jacobian_X vector tangents and @inbounds in _pb_evaluate_pbAA!.
5. **Retire the design-matrix cache** (it exhausted the storage quota at 200 GB; assembly is now minutes) and **replace svd(R)** in the factor-once solver with a Cholesky per λ or a threaded SVD – at $n = 46685$ it is the remaining post-assembly cost.
6. **Generator before basis.** Bounding the perturbations halved the error; adding short-range order and the unary/binary structures of GRACE’s “extended distillation” are the next data-side steps and cost minutes of GPU.
7. **Next accuracy lever, convexly:** fixed $\sqrt{\rho}$ Finnis–Sinclair columns with learned (VarPro) species weights inside the density gave –12% at degree 4 and transfer frozen to a disjoint split; the shape of the embedding function beyond $\sqrt{\cdot}$ does not matter. This is the like-for-like answer to GRACE-FS and should be included in any $S > 5$ run.

Sources: acejax/bench/distil/RESULTS.md; docs/findings/FINDINGS_assembly_profile.md, FINDINGS_forward_profile.md, FINDINGS_fs_spike.md, FINDINGS_varpro_spike.md, FINDINGS_codes_survey.md; plan docs/plans/jax_ace_port_plan.md (Stage 1D–E, 2B–2C).